

▼ Pytorch homework

Instructions

- Make a copy of this notebook in your own Colab and complete the questions there.
- You can add more cells if necessary. You may also add descriptions to your code, though it is not mandatory.
- Make sure the notebook can run through by *Runtime* -> *Run all*. **Keep all cell outputs** for grading.
- Submit the link of your notebook
- Please **enable editing or comments** so that you can receive feedback from TAs..

[link text](#) Install PyTorch and Skorch.

```
!pip install -q torch skorch torchvision torchtext
```

```
|████████████████████████████████████████████████████████████████████████████████| 185 kB 35.0 MB/s
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import skorch
import sklearn
import numpy as np
import matplotlib.pyplot as plt
```

▼ 1. Tensor Operations (40 points)

Tensor operations are important in deep learning models. In this part, you are required to implement some common tensor operations in PyTorch.

:### 1) Tensor squeezing, unsqueezing and viewing

Tensor squeezing, unsqueezing and viewing are important methods to change the dimension of a Tensor, and the corresponding functions are [torch.squeeze](#), [torch.unsqueeze](#) and [torch.Tensor.view](#). Please read the documents of the functions, and finish the following practice.

```

from re import X
# x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])
print(x)
x = torch.unsqueeze(x, -1) # Add two new dimensions to x by using the function torch.unsqueeze
x = torch.unsqueeze(x, 1)
print(x.shape)
x = torch.squeeze(x, 1) # Remove the two dimensions just added by using the function torch.s
x = torch.squeeze(x, 2)
print(x.shape)
x = x.view(6)
print("Tensor after resize: ", x) # x is now a two-dimensional tensor, or in other words a mat
print(x.shape)

tensor([[1., 2.],
        [3., 4.],
        [5., 6.]])
torch.Size([3, 1, 2, 1])
torch.Size([3, 2])
Tensor after resize: tensor([1., 2., 3., 4., 5., 6.])
torch.Size([6])

```

▼ 2) Tensor concatenation and stack

Tensor concatenation and stack are operations to combine small tensors into big tensors. The corresponding functions are [torch.cat](#) and [torch.stack](#). Please read the documents of the functions, and finish the following practice.

```

# x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])
x.size()
# y is a tensor with size being (3, 2)
y = torch.Tensor([[-1, -2], [-3, -4], [-5, -6]])

# Our goal is to generate a tensor z with size as (2, 3, 2), and z[0,:,:] = x, z[1,:,:] = y.
# Use torch.stack to generate such a z
print("Stack tensors ")
z = torch.stack((x,y), 0)
print("the stack of z: ", z)
print(z.shape)
# Use torch.cat and torch.unsqueeze to generate such a z
print("concatenate tensors ")
z = torch.cat((x,y), 0)
print("the torch.cat of z: ", z)
print(z.shape)
print("to unsqueeze the data")

```

```

z = torch.unsqueeze(z, 1)
print(z.shape)

Stack tensors
the stack of z: tensor([[[ 1.,  2.],
                        [ 3.,  4.],
                        [ 5.,  6.]],
                       [[-1., -2.],
                        [-3., -4.],
                        [-5., -6.]])
torch.Size([2, 3, 2])
concatenate tensors
the torch.cat of z: tensor([[ 1.,  2.],
                        [ 3.,  4.],
                        [ 5.,  6.],
                        [-1., -2.],
                        [-3., -4.],
                        [-5., -6.]])
torch.Size([6, 2])
to unsqueeze the data

```

▼ 3) Tensor expansion

Tensor expansion is to expand a tensor into a larger tensor along singleton dimensions. The corresponding functions are [torch.Tensor.expand](#) and [torch.Tensor.expand_as](#). Please read the documents of the functions, and finish the following practice.

```

# x is a tensor with size being (3)
x = torch.Tensor([1, 2, 3])
# Our goal is to generate a tensor z with size (2, 3), so that z[0,:,:] = x, z[1,:,:] = x.

# [TO DO]
# Change the size of x into (1, 3) by using torch.unsqueeze.
print("to unsqueeze the data")
x = torch.unsqueeze(x,0)
print(x.shape)

# [TO DO]
# Then expand the new tensor to the target tensor by using torch.Tensor.expand.

exp = x.expand(3, -1)
print("Expanded Tensor:", exp)
exp.size()

to unsqueeze the data
torch.Size([1, 3])
Expanded Tensor: tensor([[1., 2., 3.],
                        [1., 2., 3.]])

```

```

        [1., 2., 3.]))
torch.mean(x))

```

▼ 4) Tensor reduction in a given dimension

In deep learning, we often need to compute the mean/sum/max/min value in a given dimension of a tensor. Please read the document of [torch.mean](#), [torch.sum](#), [torch.max](#), [torch.min](#), [torch.topk](#), and finish the following practice.

```

# x is a random tensor with size being (10, 50)
x = torch.randn(10, 50)
# Compute the mean value for each row of x.
# You need to generate a tensor x_mean of size (10), and x_mean[k, :] is the mean value of th
print("mean value")
torch.mean(x)
print(x)
k=3
x_m1 = x[k]
print(x_m1)
# Compute the sum value for each row of x.
# You need to generate a tensor x_sum of size (10).
print("Sum")
print(torch.sum(x, dim=0)) # size = [1, ncol]
# Compute the max value for each row of x.
# You need to generate a tensor x_max of size (10).
print("Max")
print(torch.max(x))
# Compute the min value for each row of x.
# You need to generate a tensor x_min of size (10).
print("Min")
print(torch.min(x))
# Compute the top-5 values for each row of x.
# You need to generate a tensor x_mean of size (10, 5), and x_top[k, :] is the top-5 values o
x_m2 = torch.randn(10, 5)
print("x_mean value")
torch.mean(x_m2)
print(x_m2)
k=3
x_m2 = x_m2[k]
print(x_m2)

```

```

1.2747, -0.8570, -1.1846, 0.2547, -1.1219, 0.4294, 0.8655, -0.6636,
-0.1603, -0.9577, 1.8455, -0.0880, -0.1176, 1.4293, -1.2249, 0.5246,
1.1397, -0.6352, -0.4408, 1.4582, 0.0493, 1.4264, 1.1370, -0.9968,
-0.9809, 0.5528, -0.6590, -0.7737, -0.4931, 0.9309, -0.1715, 0.4593,
0.7562, -0.7132, 0.7413, -0.2494, 0.6437, -0.2087, -0.1612, 0.8058,
-0.2579, 0.6449],
[ 1.3088, 1.0038, -0.7025, -0.4895, 2.2443, -0.5579, 1.0837, -0.0479,
-1.0763, 0.0034, 0.9128, -0.6735, -0.2912, 0.9961, 0.2874, -0.8188,
0.2087, 2.2313, 0.8242, 1.0563, -0.0789, -0.9028, 0.2103, -1.9817,
1.2797, 0.4267, -1.0700, 1.5687, 2.1278, -0.5458, 1.8857, 0.2856,
0.3055, 0.3010, 0.3025, 0.3050, 0.3000, 0.5000, 0.3000, 0.3010]

```

```

-0.3953, -0.8918, -0.2135, 0.0359, 0.0394, -0.5803, -0.3609, 1.2043,
-0.4164, 0.5758, -2.1096, 1.0636, 0.1053, -0.1219, 0.4086, -1.0845,
-0.7320, 0.3663],
[ 0.3512, 1.4783, 1.2183, 1.8444, 0.3375, -1.2745, 1.0095, 1.0794,
-0.2481, 0.7000, 0.6696, 1.1576, -1.0951, 2.1208, 0.7194, -0.1402,
0.9188, 0.6327, 0.7467, -0.3581, 0.4409, -1.1285, 0.6091, -0.8172,
0.9734, 0.0638, 0.9751, 1.6680, 0.8366, 0.6696, -0.1859, -0.4323,
0.9329, 0.8481, 0.7597, -1.8823, 1.0147, 0.6908, 0.3829, 0.8090,
0.3277, -1.6716, -2.7289, 0.0581, -0.3082, 1.2958, 2.1448, 0.8958,
-1.1475, -0.0933],
[ 0.5948, -1.1587, -2.9344, 0.5049, -1.1358, 1.8283, 1.0256, -0.0394,
0.4751, -0.2781, 0.0601, -0.3192, 0.1615, -0.3267, 0.2448, 1.2694,
-0.2937, 1.0967, -0.9499, -0.1533, 1.0365, 0.3014, 0.5981, 0.7578,
0.7062, -0.7812, 1.6604, -1.1127, 1.1442, -0.6625, -0.0687, 0.2210,
0.1890, 1.9474, 0.8613, -1.0991, 1.2178, -1.4973, 0.0035, 0.6927,
-0.1973, 0.3492, 0.1804, -0.6620, -0.8443, -0.8268, -0.8110, -1.7725,
-1.8442, -0.0199]])
tensor([ 1.4250, -0.7966, -1.3609, 0.0917, 0.1699, -0.7673, -0.0090, 0.6805,
0.2873, -0.5812, -1.4305, -0.5231, -0.2142, -0.0890, 1.7204, 0.9599,
0.4978, -2.2911, -0.0957, -1.3734, 2.8082, 0.7507, 1.4813, -0.3913,
-0.3131, -0.0203, -0.7155, 1.0847, 0.0779, 1.0867, 1.0420, 0.1514,
-1.0449, -0.4607, 1.7630, 0.6522, 1.2171, -0.2363, 0.0608, -0.1744,
-0.5854, 0.8131, -0.3207, 0.3327, 1.3757, -0.4524, -0.5721, -1.1262,
0.8515, 1.1427])
Sum
tensor([ 3.7148, -2.3317, -8.4673, 1.3965, 2.8886, -3.0052, 5.5236, 3.8121,
3.0576, 3.2446, -3.9734, -0.3218, -2.1200, 3.3347, -1.1085, 0.7441,
3.9501, -3.2509, 4.7032, -2.1320, 0.2556, 2.8574, 2.9922, -0.2812,
4.3808, -3.5058, 2.4437, 3.2345, 0.6644, 2.2934, -1.0142, -2.4227,
-1.6704, 1.0953, 1.8204, -2.6881, 3.6140, -4.1130, -2.7999, 1.2580,
0.1238, -4.0891, -4.6261, -1.4096, -0.0595, 0.1067, 2.3059, -4.8223,
-1.4473, 3.8437])
Max
tensor(2.8082)
Min
tensor(-2.9344)
x_mean value
tensor([[ 1.0170, -0.4079, -0.0831, 0.7348, 1.2474],
[-0.9020, 0.9491, -0.5279, 0.0155, -0.3324],
[ 0.5718, -0.2169, -0.3735, 0.7183, -0.9384],
[ 1.4046, -0.0272, 0.1702, -0.7667, 0.5975],
[-1.2922, 0.2033, 1.4364, -1.4252, 0.1074],
[ 0.6342, -0.6147, -1.2083, 0.2791, 0.5005],
[-0.2909, -0.2979, -0.5955, 0.5165, -0.1222],
[ 0.0982, -1.1251, -1.5430, -1.2905, -1.0654],
[-0.3147, 0.1634, -1.0148, -0.1135, -1.5575],
[-0.5015, 1.1539, -0.6104, -0.1824, 0.7834]])
tensor([ 1.4046, -0.0272, 0.1702, -0.7667, 0.5975])

```

Install PyTorch and Skorch.

```
!pip install -q torch skorch torchvision torchtext
```

185 kB 23.7 MB/s

▼ Convolutional Neural Networks

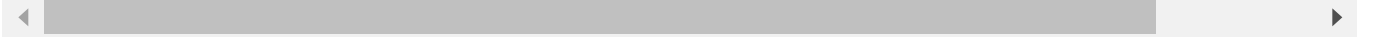
Implement a convolutional neural network for image classification on CIFAR-10 dataset.

CIFAR-10 is an image dataset of 10 categories. Each image has a size of 32x32 pixels. The following code will download the dataset, and split it into `train` and `test`. For this question, we use the default validation split generated by Skorch.

```
# importing required libraries
import torchvision
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
import skorch
from skorch.helper import predefined_split
from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d, MaxPool2d, Module, S
import numpy as np
```

```
train = torchvision.datasets.CIFAR10("./data", train=True, download=True)
test = torchvision.datasets.CIFAR10("./data", train=False, download=True)
```

```
Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to ./data/cifar-10-
100% 170498071/170498071 [00:02<00:00, 104748438.53it/s]
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified
```



The following code visualizes some samples in the dataset. You may use it to debug your model if necessary.

```
def plot(data, labels=None, num_sample=5):
    n = min(len(data), num_sample)
    for i in range(n):
        plt.subplot(1, n, i+1)
        plt.imshow(data[i], cmap="gray")
        plt.xticks([])
        plt.yticks([])
        if labels is not None:
            plt.title(labels[i])
```

```
train.labels = [train.classes[target] for target in train.targets]
plot(train.data, train.labels)
```



▼ 1) Basic CNN implementation

Consider a basic CNN model

- It has 3 convolutional layers, followed by a linear layer.
- Each convolutional layer has a kernel size of 3, a padding of 1.
- ReLU activation is applied on every hidden layer.

Please implement this model in the following section. You will need to tune the hyperparameters and fill the results in the table.

▼ a) Implement convolutional layers

Implement the initialization function and the forward function of the CNN.

```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        # implement parameter definitions here
        self.cnn_layers = Sequential(
            Conv2d(3, 4, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2), #Adding max pooling also

            # Defining another 2D convolution layer
            Conv2d(4, 4, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),

            # Defining third 2D convolution layer
            Conv2d(4, 4, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2)
        )
        self.linear_layers = Sequential(
            Linear(64, 10),
            Sigmoid())
```

```

    )

def forward(self, images):
    # implement the forward function here
    x = self.cnn_layers(images)
    x = x.view(x.size(0), -1)
    x = self.linear_layers(x)
    return x

```

▼ b) Tune hyperparameters

Train the CNN model on CIFAR-10 dataset. Tune the number of channels, optimizer, learning rate and the number of epochs for best validation accuracy.

```

# implement hyperparameters here
model = skorch.NeuralNetClassifier(CNN, criterion=torch.nn.CrossEntropyLoss,
                                   optimizer = torch.optim.SGD, lr = 0.02, max_epochs = 80)
# implement input normalization & type cast here
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
train_y = np.array(train.targets)
#train_y = torch.from_numpy(train_y)

train_y_onehot = np.zeros((train_y.size, train_y.max() + 1))
train_y_onehot[np.arange(train_y.size), train_y] = 1

train_x = np.array([[img[:, :, 0], img[:, :, 1], img[:, :, 2]] for img in train.data]).astype(float)

model.fit(torch.from_numpy(np.array(train_x, dtype = np.float32)), torch.from_numpy(train_y))

```

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.2564	0.1333	2.2229	7.6228
2	2.2117	0.1569	2.2027	7.5643
3	2.1730	0.2140	2.1594	7.6917
4	2.1352	0.2475	2.1164	7.4184
5	2.0722	0.3020	2.0506	7.7827
6	2.0272	0.2980	2.0506	7.2984
7	2.0090	0.3370	2.0130	7.3981
8	1.9965	0.3530	1.9931	7.6817
9	1.9888	0.3574	1.9945	7.8385
10	1.9812	0.3549	1.9868	7.6940
11	1.9752	0.3469	1.9929	7.5437
12	1.9696	0.3586	1.9840	7.5971
13	1.9642	0.3756	1.9667	7.7023
14	1.9619	0.3817	1.9650	7.5391
15	1.9588	0.3800	1.9606	7.8981
16	1.9554	0.3796	1.9597	7.5926
17	1.9533	0.3425	2.0062	7.5780
18	1.9512	0.3598	1.9766	8.3165

19	1.9479	0.3799	1.9641	8.3424
20	1.9461	0.3670	1.9700	7.7122
21	1.9440	0.3788	1.9521	7.8264
22	1.9423	0.3934	1.9426	7.9183
23	1.9404	0.3579	1.9905	7.6476
24	1.9393	0.3814	1.9571	7.7644
25	1.9364	0.3780	1.9508	7.5503
26	1.9387	0.3446	1.9856	7.7931
27	1.9363	0.3906	1.9464	7.7912
28	1.9372	0.3764	1.9508	7.7820
29	1.9351	0.3921	1.9362	7.6377
30	1.9356	0.3687	1.9623	7.5101
31	1.9345	0.3806	1.9463	7.7982
32	1.9342	0.3808	1.9479	7.8631
33	1.9326	0.3739	1.9534	7.4842
34	1.9331	0.3849	1.9476	7.6901
35	1.9303	0.3918	1.9329	7.7818
36	1.9319	0.3743	1.9486	7.7786
37	1.9302	0.3704	1.9485	7.6469
38	1.9323	0.3787	1.9399	7.8281
39	1.9295	0.3853	1.9358	7.7532
40	1.9280	0.3635	1.9554	7.5893
41	1.9273	0.3751	1.9494	7.6380
42	1.9289	0.3937	1.9327	7.7828
43	1.9289	0.3324	1.9669	7.7377
44	1.9292	0.3607	1.9508	7.5954
45	1.9298	0.3328	1.9712	7.7538
46	1.9299	0.3709	1.9409	7.8036
47	1.9310	0.3483	1.9646	7.7836
48	1.9304	0.3610	1.9529	7.7242
49	1.9280	0.3776	1.9306	7.7271
50	1.9300	0.3577	1.9457	7.7726
51	1.9292	0.3389	1.9560	7.6273
52	1.9289	0.3482	1.9424	7.4246
53	1.9284	0.3339	1.9471	7.8989
54	1.9290	0.3600	1.9351	7.5079
55	1.9296	0.3497	1.9406	7.5861
56	1.9295	0.3761	1.9777	7.7060

Validation loss obtained using SGD optimizer for 50 epochs is 1.9457 Validation loss obtained using SGD optimizer for 80 epochs is 1.9392

```
# implement hyperparameters here
model = skorch.NeuralNetClassifier(CNN, criterion=torch.nn.CrossEntropyLoss,
                                   optimizer = torch.optim.Adam, lr = 0.2, max_epochs = 50)

# implement input normalization & type cast here
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
train_y = np.array(train.targets)
#train_y = torch.from_numpy(train_y)

train_y_onehot = np.zeros((train_y.size, train_y.max() + 1))
train_y_onehot[np.arange(train_y.size), train_y] = 1
```

```
train_x = np.array([[img[:, :, 0], img[:, :, 1], img[:, :, 2]] for img in train.data]).astype(float)

model.fit(torch.from_numpy(np.array(train_x, dtype = np.float32)), torch.from_numpy(train_y))
```

epoch	train_loss	valid_acc	valid_loss	dur
1	2.3616	0.1000	2.3612	8.6093
2	2.3612	0.1000	2.3612	9.9424
3	2.3612	0.1000	2.3612	8.3205
4	2.3612	0.1000	2.3612	7.8114
5	2.3612	0.1000	2.3612	7.7383
6	2.3612	0.1000	2.3612	7.7371
7	2.3612	0.1000	2.3612	7.6811
8	2.3612	0.1000	2.3612	7.6484
9	2.3612	0.1000	2.3612	7.7050
10	2.3612	0.1000	2.3612	7.9149
11	2.3612	0.1000	2.3612	7.7035
12	2.3612	0.1000	2.3612	8.2284
13	2.3612	0.1000	2.3612	7.7088
14	2.3612	0.1000	2.3612	7.6084
15	2.3612	0.1000	2.3612	7.7143
16	2.3612	0.1000	2.3612	7.8593
17	2.3612	0.1000	2.3612	7.6511
18	2.3612	0.1000	2.3612	7.8716
19	2.3612	0.1000	2.3612	10.9168
20	2.3612	0.1000	2.3612	8.3139
21	2.3612	0.1000	2.3612	8.2782
22	2.3612	0.1000	2.3612	9.4205
23	2.3612	0.1000	2.3612	8.5033
24	2.3612	0.1000	2.3612	7.8255
25	2.3612	0.1000	2.3612	8.9048
26	2.3612	0.1000	2.3612	9.9808
27	2.3612	0.1000	2.3612	9.4702
28	2.3612	0.1000	2.3612	9.6823
29	2.3612	0.1000	2.3612	10.1683
30	2.3612	0.1000	2.3612	9.4410
31	2.3612	0.1000	2.3612	10.5238
32	2.3612	0.1000	2.3612	12.1579
33	2.3612	0.1000	2.3612	8.7332
34	2.3612	0.1000	2.3612	8.1706
35	2.3612	0.1000	2.3612	8.3132
36	2.3612	0.1000	2.3612	8.1552
37	2.3612	0.1000	2.3612	8.2883
38	2.3612	0.1000	2.3612	8.0469
39	2.3612	0.1000	2.3612	8.2140
40	2.3612	0.1000	2.3612	8.2633
41	2.3612	0.1000	2.3612	8.2459
42	2.3612	0.1000	2.3612	8.1746
43	2.3612	0.1000	2.3612	8.2453
44	2.3612	0.1000	2.3612	8.1371
45	2.3612	0.1000	2.3612	8.7439
46	2.3612	0.1000	2.3612	8.3979
47	2.3612	0.1000	2.3612	8.2403
48	2.3612	0.1000	2.3612	8.2722
49	2.3612	0.1000	2.3612	8.1669
50	2.3612	0.1000	2.3612	8.1890

```
<class 'skorch.classifier.NeuralNetClassifier'>[initialized](
  module_=CNN(
    (cnn_layers): Sequential(
      (0): Conv2d(3, 4, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU(inplace=True)
      (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

Adam optimizer could achieve 2.3612

Write down **validation accuracy** of your model under different hyperparameter settings. Note the validation set is automatically split by Skorch during `model.fit()`.

Hint: You may need more epochs for SGD than Adam.

#channel for each layer \ optimizer	SGD	Adam
(128, 128, 128)		
(256, 256, 256)		
(512, 512, 512)		

▼ c) Use larger CNN model

Indented block

Add more Convolution/BatchNorm/Pooling/DropOut/Linear layers to improve the accuracy.

```
class CNNLarge(nn.Module):
    def __init__(self):
        super(CNNLarge, self).__init__()
        # implement parameter definitions here
        self.cnn_layers = Sequential(
            Conv2d(3, 4, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2), #Adding max pooling also

            # Defining another 2D convolution layer
            Conv2d(4, 4, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),

            # Defining third 2D convolution layer
            Conv2d(4, 8, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),

            # Defining fourth 2D convolution layer
            Conv2d(8, 16, kernel_size=3, stride=1, padding=1),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),
```

```

        # Defining fifth 2D convolution layer
        Conv2d(16, 32, kernel_size=3, stride=1, padding=1),
        ReLU(inplace=True),
        MaxPool2d(kernel_size=2, stride=2)

    )
    self.linear_layers = Sequential(
        Linear(32, 10)
    )

def forward(self, images):
    # implement the forward function here
    x = self.cnn_layers(images)
    x = x.view(x.size(0), -1)
    x = self.linear_layers(x)
    return x

# implement hyperparameters here
model = skorch.NeuralNetClassifier(CNNLarge, criterion=torch.nn.CrossEntropyLoss,
                                  optimizer = torch.optim.SGD, lr = 0.02, max_epochs = 50)
# implement input normalization & type cast here
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
train_y = np.array(train.targets)
#train_y = torch.from_numpy(train_y)

train_y_onehot = np.zeros((train_y.size, train_y.max() + 1))
train_y_onehot[np.arange(train_y.size), train_y] = 1

train_x = np.array([[img[:, :, 0], img[:, :, 1], img[:, :, 2]] for img in train.data]).astype(float)

model.fit(torch.from_numpy(np.array(train_x, dtype = np.float32)), torch.from_numpy(train_y))

```

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2608	0.2042	2.1160	8.4619
2	1.9980	0.2766	1.9392	8.4601
3	1.8505	0.3473	1.7859	8.2478
4	1.7526	0.3628	1.7747	8.3604
5	1.6938	0.3786	1.6956	8.5037
6	1.6556	0.3923	1.6679	8.3975
7	1.6272	0.4118	1.6227	8.4142
8	1.6061	0.4194	1.5975	8.5926
9	1.5875	0.4146	1.5976	9.4106
10	1.5706	0.4283	1.5837	8.5247
11	1.5544	0.4286	1.5701	8.5465
12	1.5415	0.3656	1.7803	8.6729
13	1.5263	0.4024	1.6732	11.3805
14	1.5180	0.4232	1.5944	11.6643
15	1.5050	0.4184	1.6173	11.9321
16	1.4906	0.4324	1.5705	10.6959
17	1.4830	0.4374	1.5637	9.7695
18	1.4742	0.4401	1.5571	9.1306

19	1.4693	0.4411	1.5321	9.2395
20	1.4623	0.4429	1.5604	8.4874
21	1.4520	0.4378	1.5846	8.6886
22	1.4475	0.4306	1.6232	8.7709
23	1.4356	0.4464	1.5482	8.8958
24	1.4340	0.4372	1.5922	8.5705
25	1.4292	0.4603	1.5276	8.5893
26	1.4172	0.4540	1.5488	9.2090
27	1.4140	0.4513	1.5469	8.7739
28	1.4109	0.4559	1.5338	8.7279
29	1.4060	0.4597	1.5139	8.6909
30	1.3977	0.4374	1.5963	8.5998
31	1.3965	0.4642	1.5009	8.5772
32	1.3926	0.4575	1.5342	8.5082
33	1.3864	0.4620	1.5183	8.4655
34	1.3873	0.4774	1.4784	9.7116
35	1.3794	0.4670	1.4861	8.8878
36	1.3757	0.4695	1.4838	11.4243
37	1.3697	0.4779	1.4915	8.7505
38	1.3688	0.4635	1.5149	8.7212
39	1.3695	0.4661	1.4914	8.5417
40	1.3608	0.4493	1.5576	8.5625
41	1.3597	0.4741	1.4798	8.6888
42	1.3676	0.4820	1.4628	8.6457
43	1.3596	0.4723	1.4801	8.5798
44	1.3559	0.4724	1.4850	8.6249
45	1.3505	0.4763	1.4725	8.6494
46	1.3485	0.4780	1.4737	8.6372
47	1.3464	0.4834	1.4576	8.6884
48	1.3438	0.4827	1.4550	8.7124
49	1.3405	0.4620	1.5084	8.6828
50	1.3427	0.4775	1.4682	8.5796

```
<class 'skorch.classifier.NeuralNetClassifier'>[initialized](
  module_=CNNLarge(
    (cnn_layers): Sequential(
      (0): Conv2d(3, 4, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): ReLU(inplace=True)
      (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

Using a larger architecture, I could obtain much better loss of 1.4682. This clearly proves that this problem requires complexity and using bigger models can give better results

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 5:48 PM

